

yd2459_HW3

Yihang Ding

10/21/2019

P3 Use NewtonRaphson algorithm from Section 4.4.1, pp. 120-121, [ESL] book and perform 10 iterations.
Hint: Use library(matlib) for calculating matrix inverses.

```
library(matlab)

##
## Attaching package: 'matlab'

## The following object is masked from 'package:stats':
##
##      reshape

## The following objects are masked from 'package:utils':
##
##      find, fix

## The following object is masked from 'package:base':
##
##      sum

x = c(0.0, 0.2, 0.4, 0.6, 0.8, 1.0)
y = c(0, 0, 0, 1, 0, 1)
first = function(beta){
  f1 = sum((exp(beta[1]+beta[2]*x)/(1+exp(beta[1]+beta[2]*x)))-y)
  f2 = sum(((exp(beta[1]+beta[2]*x)/(1+exp(beta[1]+beta[2]*x)))-y)*x)
  matrix(c(f1, f2), ncol=1, nrow=2)
}

second = function(beta){
  f1 = sum(exp(beta[1]+beta[2]*x)/(1+exp(beta[1]+beta[2]*x))^2)
  f2 = sum(exp(beta[1]+beta[2]*x)/(1+exp(beta[1]+beta[2]*x))^2*x)
  f3 = sum(exp(beta[1]+beta[2]*x)/(1+exp(beta[1]+beta[2]*x))^2*x)
  f4 = sum(exp(beta[1]+beta[2]*x)/(1+exp(beta[1]+beta[2]*x))^2*x*x)
  matrix(c(f1, f2, f3, f4), ncol = 2, nrow = 2)
}

beta = c(0, 0)
sequence = seq(1,10)
b1 = seq(1,10)
b2 = seq(1,10)
for(val in sequence){
  # update: new = old - second^-1 * first
  beta = beta - solve(second(beta))%*%first(beta)
  cat("Update ", val, ",beta0: ",beta[1], ", beta1:", beta[2],"\n")
  b1[val] = beta[1]
  b2[val] = beta[2]
}
```

```
## Update 1 ,beta0: -2.380952 , beta1: 3.428571
## Update 2 ,beta0: -3.522775 , beta1: 4.966947
## Update 3 ,beta0: -4.022333 , beta1: 5.624766
## Update 4 ,beta0: -4.096585 , beta1: 5.721513
## Update 5 ,beta0: -4.09797 , beta1: 5.723308
## Update 6 ,beta0: -4.09797 , beta1: 5.723309
## Update 7 ,beta0: -4.09797 , beta1: 5.723309
## Update 8 ,beta0: -4.09797 , beta1: 5.723309
## Update 9 ,beta0: -4.09797 , beta1: 5.723309
## Update 10 ,beta0: -4.09797 , beta1: 5.723309
```

```
b1
```

```
## [1] -2.380952 -3.522775 -4.022333 -4.096585 -4.097970 -4.097970 -4.097970
## [8] -4.097970 -4.097970 -4.097970
```

```
b2
```

```
## [1] 3.428571 4.966947 5.624766 5.721513 5.723308 5.723309 5.723309
## [8] 5.723309 5.723309 5.723309
```

```
select = glm(y~x,family = "binomial")
select$coefficients[1]
```

```
## (Intercept)
## -4.09797
```

```
select$coefficients[2]
```

```
## x
## 5.723309
```

Answer: $\beta_0 = -4.098$, $\beta_1 = 5.723$

P7 (a) First run `set.seed(1000)`, and then create a training set containing a random sample of 800 observations, and a test set containing the remaining observations.

```
library(ISLR)
set.seed(1000)
dim(OJ)
```

```
## [1] 1070 18
```

```
sample = sample.int(n = dim(OJ)[1], size = 800, replace = F)
train = OJ[sample, ]
test = OJ[-sample, ]
dim(train)
```

```
## [1] 800 18
```

```
dim(test)
```

```
## [1] 270 18
```

- (b) Fit a tree to the training data, with Purchase as the response and the other variables as predictors. Use the `summary()` function to produce summary statistics about the tree, and describe the results obtained. What is the training error rate? How many terminal nodes does the tree have?

```
library(tree)
fit = tree(Purchase~., data = train)
summary(fit)
```

```
##
## Classification tree:
## tree(formula = Purchase ~ ., data = train)
## Variables actually used in tree construction:
## [1] "LoyalCH"      "PriceDiff"    "SalePriceMM"
## Number of terminal nodes: 8
## Residual mean deviance: 0.7486 = 592.9 / 792
## Misclassification error rate: 0.16 = 128 / 800
```

Training error rate is 0.16, the tree has 8 terminal nodes.

- (c) Type in the name of the tree object in order to get a detailed text output. Pick one of the terminal nodes, and interpret the information displayed.


```
fit
```

```
## node), split, n, deviance, yval, (yprob)
##      * denotes terminal node
##
## 1) root 800 1066.00 CH ( 0.61500 0.38500 )
##    2) LoyalCH < 0.5036 353 422.60 MM ( 0.28612 0.71388 )
##      4) LoyalCH < 0.276142 170 131.00 MM ( 0.12941 0.87059 )
##        8) LoyalCH < 0.035047 57 10.07 MM ( 0.01754 0.98246 ) *
##        9) LoyalCH > 0.035047 113 108.50 MM ( 0.18584 0.81416 ) *
##      5) LoyalCH > 0.276142 183 250.30 MM ( 0.43169 0.56831 )
##        10) PriceDiff < 0.05 78 79.16 MM ( 0.20513 0.79487 ) *
##        11) PriceDiff > 0.05 105 141.30 CH ( 0.60000 0.40000 ) *
##    3) LoyalCH > 0.5036 447 337.30 CH ( 0.87472 0.12528 )
##      6) LoyalCH < 0.764572 187 206.40 CH ( 0.75936 0.24064 )
##        12) SalePriceMM < 2.125 120 156.60 CH ( 0.64167 0.35833 )
##          24) PriceDiff < -0.35 16 17.99 MM ( 0.25000 0.75000 ) *
##          25) PriceDiff > -0.35 104 126.70 CH ( 0.70192 0.29808 ) *
##        13) SalePriceMM > 2.125 67 17.99 CH ( 0.97015 0.02985 ) *
##      7) LoyalCH > 0.764572 260 91.11 CH ( 0.95769 0.04231 ) *
```

Take "9) LoyalCH > 0.035047 113 108.50 MM (0.18584 0.81416) *" for example, it uses LoyalCH with threshold 0.35047 as split feature, and get 113 samples. With deviance 108.5 and prediction value as MM.

- (d) (2pt) Create a plot of the tree, and interpret the results.

```
plot(fit, main='OJ train Decision Tree')
text(fit)
```



- (e) (3pt) Predict the response on the test data, and produce a confusion matrix comparing the test labels to the predicted test labels. What is the test error rate?

```
pred = predict(fit, test, type = 'class')
summary(pred)
```

```
## CH MM
## 188 82
```

```
table(pred, test$Purchase)
```

```
##
## pred CH MM
## CH 150 38
## MM 11 71
```

```
mean(pred != test$Purchase)
```

```
## [1] 0.1814815
```

Test error rate is 0.18.

(f) (2pt) Apply the `cv.tree()` function to the training set in order to determine the optimal tree size.


```
cv_fit = cv.tree(fit, FUN=prune.misclass)
cv_fit

## $size
## [1] 8 7 4 2 1
##
## $dev
## [1] 142 142 143 164 308
##
## $k
## [1]      -Inf    0.000000    2.666667   10.500000  151.000000
##
## $method
## [1] "misclass"
##
## attr(,"class")
## [1] "prune"          "tree.sequence"
```

Optimal tree size is 4.

(g) (3pt) Produce a plot with tree size on the x-axis and cross-validated classification error rate on the y-axis.

```
plot(x=cv_fit$size, y=cv_fit$k, type='b', xlab='Tree Size', ylab='CV err rate')
```



- (h) (1pt) Which tree size corresponds to the lowest cross-validated classification error rate? Tree with size 4 has the lowest CV classification error rate.
- (i) (3pt) Produce a pruned tree corresponding to the optimal tree size obtained using cross-validation. If cross-validation does not lead to selection of a pruned tree, then create a pruned tree with five terminal nodes.

```
prune = prune.misclass(fit, best=4)
summary(prune)
```

```
##
## Classification tree:
## snip.tree(tree = fit, nodes = 4:3)
## Variables actually used in tree construction:
## [1] "LoyalCH" "PriceDiff"
## Number of terminal nodes: 4
## Residual mean deviance: 0.8653 = 688.8 / 796
## Misclassification error rate: 0.17 = 136 / 800
```

- (j) (2pt) Compare the training error rate between the pruned and unpruned tree. Which is higher? Pruned tree has higher train error rate.
- (k) (2pt) Compare the test error rates between the pruned and unpruned trees. Which is higher?

```
prune_pred = predict(prune, test, type = 'class')  
mean(prune_pred != test$Purchase)
```

```
## [1] 0.2037037
```

Pruned tree has higher test error rate.